Samarth Garg

24 Lowther Road, George Town, Prayagraj, Uttar Pradesh, India | samarthgarg8099@gmail.com

+91 7007 451 423 | LinkedIn | GitHub | LeetCode | Kaggle | 22ucs175@lnmiit.ac.in

Education

LNM Institute of Information Technology, Jaipur, Rajasthan, B.Tech (Hons.) in Computer Science and Engineering with Specialization in Artificial Intelligence and Data Science — Oct 2022 – May 2026

Technical Skills

Languages: Python, Java, SQL
Libraries: Pytest, NumPy, Pandas, Seaborn, NLTK, Scikit-learn, Hugging Face Transformers
Databases and Vector Stores: MySQL, MongoDB, PostgreSQL, FAISS, ChromaDB
Tools: Git, MLflow, Docker, GitHub Actions, Azure, Kafka, Postman, Prometheus
Frameworks: TensorFlow, Flask, FastAPI, Streamlit, LangChain, LangGraph, LlamaIndex, PySpark, Prefect
Coursework: Machine Learning, Deep Learning, Big Data Analytics, Generative AI, Introduction to LLMs

Work Experience

Software Engineer Intern | Appaleon Consulting — Sep 2026 – Present

- Architected multi-tenant FastAPI backend **(1,025 LOC)** with JWT auth, **6+** job-management APIs, and zero-disk transient resume parsing (PDF/DOCX).
- Instituted automated Pytest infrastructure, authoring **130+** unit/integration tests with mocked database sessions to hit **68%** code coverage.

ML Engineer Intern | Manthhan — Jan 2026 – Jun 2026

- Resolved **GitHub Actions** CI failures during Python 3.12 to 3.14 migration by optimizing PR check workflows and configuring **3 GB** swap memory to overcome **AWS CodeBuild** storage constraints, ensuring stable package installation and test execution.
- Decreased **Docker** image size from **3.23 GB** to **1.73 GB (47% reduction)**, reducing **AWS Elastic Container Service** overhead.
- Improved document classification accuracy from **32%** to **95%** by optimizing regex-based pattern extraction on OCR-processed text, increasing the robustness of rule-based classification.
- Engineered an **LLM-driven** document processing system that reduced manual reading effort by **70%**.

AI Engineer Intern | YourAllyStack — Sep 2025 – Dec 2025

- Built an end-to-end interview evaluation pipeline that converted S3 video interviews into transcripts and structured Q&A pairs **(+/-1 error margin)** with automated student and rubric-based admin evaluation.
- Optimized Q&A pair extraction from **1-hour** transcripts to under **1 minute** by implementing strategic **text chunking**, enabling efficient processing for both manual and AI-conducted interview evaluations.
- Enabled scalable rubric-driven LLM evaluation by parsing Excel rubrics into dynamic scoring indicators across interview sessions, reducing token cost by **60%** and latency by **80%**.
- Integrated **PostgreSQL** storage for transcripts, scores, and feedback across **500+** interview sessions, enabling structured retrieval and audit trails.

Projects

AI-Powered Resume Screening Assistant | GitHub — Jul 2026 – Aug 2026

- Built a resume screening pipeline that processes zip uploads (files in MinIO, metadata in MySQL), ranks candidates in **under 2 seconds** via hybrid retrieval (FAISS + BM25) and Cross-Encoder reranking, achieving **100% Precision at 5 and NDCG at 5**, and scores shortlisted profiles concurrently via LLaMA 3.3 (Groq).
- **Tools Used**: Python, LlamaIndex, LLaMA 3.3 (Groq), FastAPI, Streamlit, FAISS, MinIO, MySQL, Docker, spaCy

Achievements & Certifications

Semifinalist, Rebel Foods' EatSure Cloud 2024 Hackathon — Sep 2024 – Oct 2024